

STRATEGIC PROPOSAL REQUEST

Wareed Sovereign Agentic AI Operations Center

AI-First Enterprise Operating Model for Healthcare Operations

Issued by

Wareed Medical Laboratories

Kingdom of Saudi Arabia

CONFIDENTIAL | SEPTEMBER 2026

Design Principle

AI-first, dashboard-second, room-optional. The physical command room is repositioned as an optional experience layer; the core investment is sovereign agentic AI, trusted data, integrations, governed automation, and measurable operational impact.

Document Control

Document	Strategic Proposal Request - Wareed Sovereign Agentic AI Operations Center
Purpose	Invite an AI transformation partner to design and implement a phased sovereign agentic operating model.
Basis	Reframed from the approved Wareed Business Operations Center (WBOC) RFP - Volumes I, II and III.
Issue	Draft for AI Company Engagement
Date	September 2026
Classification	Confidential

Source Positioning

The original WBOC RFP already defines the program as a business transformation initiative - not merely a conventional control room - and calls for live data, integrated systems, predictive insight, AI-enabled capabilities, proactive alerts, and enterprise-wide operational coordination. This proposal preserves those business objectives while replacing the room-centric delivery model with an agentic AI operating model.

Important Scope Principle

Where this proposal introduces Sovereign AI, agent autonomy levels, an AI control plane, and specific agent roles, these are Wareed strategic additions derived from the original objectives; they are not represented as verbatim requirements from the legacy RFP.

Contents

- 1. Executive Summary
- 2. Why the WBOC Should Become Agentic AI
- 3. Target Agentic Operating Model
- 4. Agent Architecture at a Glance
- 5. Detailed Agent Roles
- 6. Human-in-the-Loop Authority Model
- 7. Sovereign AI, Data and Cybersecurity
- 8. Integration and Data Foundation
- 9. Phased Implementation Roadmap
- 10. Proof of Concept
- 11. Acceptance and Success Measures
- 12. Vendor Response Requirements
- 13. Recommended Evaluation Model
- 14. Commercial Structure
- Appendix A - RFP-to-Agent Mapping

1. Executive Summary

Wareed Medical Laboratories is seeking an AI transformation partner to evolve the approved Wareed Business Operations Center concept into a Sovereign Agentic AI Operations Center. The new model will not depend on teams continuously watching screens. Instead, domain-specific AI agents will monitor enterprise operations, correlate data, detect exceptions, identify likely root causes, recommend or initiate approved corrective workflows, track resolution, and escalate only decisions that require human or executive judgment.

The original WBOC business requirements already call for enterprise visibility, faster decision-making, reduced manual follow-up, proactive alerts, measurable KPIs, root-cause analysis, corrective action tracking, predictive intelligence, and support for future AI capabilities. The new approach operationalizes those goals through an agentic architecture rather than treating AI as a future add-on.

Target Outcome

A digital operating model in which Wareed is continuously observed by trusted AI agents, operational exceptions are converted into governed action, and the CEO receives concise decision intelligence rather than raw dashboards.

Strategic Outcomes

- Reduce manual operational follow-up across branches, laboratory, Home Visit, courier/logistics, call center, quality, growth, IT, HR/workforce and executive leadership.
- Convert live and near-real-time data into prioritized operational intelligence and recommended action.
- Create measurable accountability through automated ownership, SLA tracking, escalation and closure evidence.
- Improve branch performance, laboratory TAT, Home Visit productivity, customer experience, marketing effectiveness, quality compliance and launch readiness.
- Provide the CEO and executive leadership with a 30-second view plus an AI-generated daily decision brief.
- Build a sovereign, auditable and scalable AI foundation that can expand to new branches, services and future agent use cases.

What Changes from the Legacy Command Room

Legacy Emphasis	Agentic AI Replacement	Commercial Treatment
Video wall / large screens	Executive dashboards remain available, but decision intelligence is pushed proactively by AI agents.	Optional / phased
Operator consoles	Human supervisors focus on exceptions and approvals, not continuous passive monitoring.	Reduced / role-based
Physical room infrastructure	Digital agentic operating environment becomes the primary platform.	Optional experience layer
Dashboards as primary interface	Agents monitor dashboards/data sources continuously and surface decisions, alerts and actions.	Core AI platform
AI as future capability	AI becomes the central operating model from Phase 1 onward.	Core mandatory scope

2. Why the WBOC Should Become Agentic AI

The original RFP identifies the core problem correctly: fragmented information, delayed decision-making, manual follow-up, reactive issue management, limited data interoperability, and the need for predictive operational intelligence. A traditional control room can visualize these problems; an agentic platform can continuously work on them.

Transformation Logic

Agentic Capability	Business Meaning
See	Connect operational data and create trusted visibility.
Understand	Correlate signals across departments and identify anomalies or root causes.
Recommend	Generate evidence-based corrective actions and decision options.
Assign	Create approved tasks or tickets and route them to accountable owners.
Escalate	Automatically escalate unresolved, repeated or high-severity issues.
Learn	Measure outcomes and improve thresholds, recommendations and playbooks over time.

RFP Foundation Preserved

- Enterprise-wide operational visibility and coordinated decision-making.
- Real-time or near-real-time dashboards, drill-down, KPI thresholds and critical alerts.
- Predictive insights, AI-driven recommendations and recurring issue detection.
- Integration with LIS, ERP, CRM, HR/attendance, GPS/fleet, CCTV/VMS, call center and customer-experience platforms.
- Data quality, master data, auditability, cybersecurity, RBAC, encryption, least privilege, backup and recovery.
- Testing, UAT, integration testing, cybersecurity testing, operational readiness, training, SLAs and support.

3. Target Agentic Operating Model

The recommended design is a multi-agent system with a central Agentic Orchestrator, domain agents, and shared control services. Each agent has a constrained mission, approved data access, explicit action boundaries, an audit trail, and escalation rules.

Architecture Layers

Layer 1 - Sovereign Data & Integration

Operational data layer, APIs/connectors, event bus, master data, data quality, identity, security and audit logs.

Layer 2 - Agentic Control Plane

Agent registry, orchestration, memory/context boundaries, policy engine, approval workflows, tool permissions, logging and observability.

Layer 3 - Domain Agents

Branch, Lab, Home Visit, Logistics, CX, Marketing, Product Launch, Quality, Workforce, Voice of Branch, IT/Integration.

Layer 4 - Shared Agentic Services

Data Quality & Governance, Alert/Ticketing/Corrective Action, AI Vision Event service.

Layer 5 - Executive Experience

CEO Briefing Agent, role-based dashboards, mobile/web alerts, drill-down and executive decision queue.

Core Principle

Human accountability remains intact

AI may monitor, analyze, recommend, assign and escalate within approved boundaries. Clinical decisions, pricing changes, campaign launches, sensitive patient communications and other high-risk actions remain subject to defined human approval.

4. Agent Architecture at a Glance

ID	Agent / Service	Primary Role	Phase	Type
A0	Agentic Orchestrator & Operations Control	Cross-enterprise orchestration, prioritization, policy, agent coordination and escalation.	Foundation	Domain Agent
A1	Executive Intelligence & CEO Briefing Agent	30-second enterprise view, daily decision brief, risk/opportunity prioritization.	MVP	Domain Agent
A2	Branch Operations Agent	Branch status, volume, wait time, revenue, staffing, complaints, risk and corrective action.	MVP	Domain Agent
A3	Main Laboratory Operations & TAT Agent	Sample flow, TAT, QC, analyzers, rejections, reagent risk and bottlenecks.	MVP	Domain Agent
A4	Home Visit Operations Agent	Visits, technicians, routes, delays, ETA, productivity, capacity and bottlenecks.	MVP	Domain Agent
A5	Customer Experience & Call Center Agent	Complaints, satisfaction, contact-center indicators, unresolved cases and recurring themes.	MVP	Domain Agent
A6	Marketing & Growth Intelligence Agent	Campaigns, acquisition, offers, packages, branch revenue, channel response and commercial actions.	MVP	Domain Agent
A7	Logistics & Courier Agent	Courier/logistics status, route/fleet signals and service exceptions.	Extended	Domain Agent
A8	Product Launch & Commercial Activation Agent	Readiness, training, materials, branch adoption, launch delays and early performance.	Extended	Domain Agent
A9	Quality & Compliance Agent	Audits, complaints, CAPA, hygiene, PPE, uniform, cleaning, recurring non-compliance.	Extended	Domain Agent
A10	Workforce & Attendance Agent	Attendance, staffing availability and workforce-related operating risk.	Extended	Domain Agent
A11	Voice of Branch / Field Intelligence Agent	Classifies field submissions, competitor observations, shortages, training needs and ideas.	Extended	Domain Agent
A12	IT & Integration Health Agent	System/interface health, API failures, refresh status, retries, logs and availability.	Foundation/Extended	Domain Agent
S1	Data Quality & Governance Agentic Service	Data validation, ownership, MDM, quality exceptions, retention and auditability.	Foundation	Shared Service
S2	Alert, Ticketing & Corrective Action Agentic Service	Severity, deduplication, owner assignment, SLA, escalation, closure and analytics.	Foundation	Shared Service
S3	AI Vision Event Agentic Service	Event-driven camera analytics with privacy controls and	Extended	Shared Service

ID	Agent / Service	Primary Role	Phase	Type
		false-positive management.		

5. Detailed Agent Roles

Each role below is directly mapped to the legacy RFP domains. Where the legacy RFP contains only limited requirements (for example, courier/logistics and workforce), the proposed agent is deliberately constrained to that supported scope unless Wareed separately expands it.

A0 | Foundation

Agentic Orchestrator & Operations Control

Mission

Act as the central control plane that coordinates domain agents, applies policy, prioritizes enterprise exceptions and routes decisions to the right human or system workflow.

What the Agent Monitors

- All agent-generated alerts and recommendations
- Enterprise severity and business impact
- Open tickets, owners, SLAs and overdue actions
- Cross-domain dependencies and repeated issues
- Human approval queues and escalation status

How the Agent Reasons

- Correlates issues across branch, lab, Home Visit, CX, quality and growth
- Prioritizes issues by severity, urgency and business impact
- Suppresses duplicate or low-value alerts
- Determines which agent or owner should act next

Actions the Agent May Perform

- Create/route approved tickets or tasks
- Assign accountable owner according to policy
- Escalate when response targets are exceeded
- Maintain complete action and decision history
- Feed resolved outcomes back into recurring-issue analytics

When It Escalates

- Executive-severity events
- Repeated unresolved issues
- Cross-functional blockers
- Actions outside AI authority
- Conflicting recommendations from domain agents

Core KPIs

- Alert-to-owner time
- SLA compliance
- Escalation accuracy
- Duplicate alert reduction
- Corrective action closure rate

Primary Data / Tool Connections

- Alert engine
- Ticketing/work management
- All domain agents
- Role/identity service
- Enterprise event logs

Human / Policy Boundary

No independent high-risk business action. The Orchestrator executes only pre-approved workflow actions and routes all other actions for human approval.

Legacy RFP basis: BR-004 to BR-010; FR-INT-001 to FR-INT-010; AR-001 to AR-012.

A1 | MVP

Executive Intelligence & CEO Briefing Agent

Mission

Convert enterprise operational data into a concise decision-oriented view for the CEO and executive leadership.

What the Agent Monitors

- Revenue and customer volume
- Branch performance and risk
- Home Visit performance
- Laboratory TAT and QC
- Quality and customer satisfaction
- Workforce availability
- Marketing/growth performance
- Critical operational and business alerts

How the Agent Reasons

- Identifies the few issues that materially require executive attention
- Correlates operational risk with financial or customer impact
- Summarizes likely root cause and options
- Separates information items from true decision items

Actions the Agent May Perform

- Generate daily CEO brief
- Create executive decision queue
- Provide 30-second enterprise summary
- Allow drill-down to region, city, branch, service, issue or root cause
- Generate weekly executive trend summary

When It Escalates

- Critical safety/quality or continuity risks
- Material revenue decline
- Enterprise-wide system failure
- Persistent TAT or Home Visit disruption
- High-impact reputation/CX issue

Core KPIs

- Brief relevance/precision
- Decision response time
- Number of avoidable executive alerts
- Time from event to executive awareness
- Action closure after executive decision

Primary Data / Tool Connections

- All domain agents
- Executive KPI layer
- Finance/revenue feeds
- Alert service
- Historical trend store

Human / Policy Boundary

Summarizes and recommends. Executive decisions remain with authorized Wareed leaders.

Legacy RFP basis: SR-001, SR-009; FR-EXE-001 to FR-EXE-015; FR-INT-010.

A2 | MVP

Branch Operations Agent

Mission

Continuously monitor branch operating health and convert branch-level performance deviations into prioritized corrective action.

What the Agent Monitors

- Open/closed status
- Customer volume and waiting time
- Daily branch revenue and target performance
- Staff attendance/availability
- Customer satisfaction and complaints
- Camera/VMS operational status where available
- Normal / Warning / Critical status
- Corrective action lifecycle

How the Agent Reasons

- Detects abnormal decline versus branch baseline or peer branches
- Ranks branches by risk and performance
- Correlates staffing, waiting time, complaints and revenue
- Identifies recurring underperformance patterns

Actions the Agent May Perform

- Create branch risk alert
- Recommend corrective action
- Assign branch/cluster owner
- Track action to closure
- Produce regional/city comparisons and branch ranking

When It Escalates

- Critical branch service disruption
- Repeated underperformance
- Unresolved complaints with business impact
- Staffing gap affecting service
- Revenue decline beyond approved threshold

Core KPIs

- Branch risk detection lead time
- Waiting time reduction
- Revenue recovery after intervention
- Corrective action closure
- Complaint recurrence

Primary Data / Tool Connections

- Branch systems
- LIS/registration where relevant
- POS/finance
- HR/attendance
- CX platform
- VMS/camera status

Human / Policy Boundary

May recommend and open tasks. Branch closure, pricing, staffing changes or customer communications require authorized approval.

Legacy RFP basis: SR-002; FR-BR-001 to FR-BR-015; FR-OPS-002; FR-EXE-005.

A3 | MVP

Main Laboratory Operations & TAT Agent

Mission

Protect laboratory service performance by continuously monitoring sample flow, analyzer status, QC, TAT and operational bottlenecks.

What the Agent Monitors

- Sample receipt through result release
- Pending and delayed samples
- STAT samples
- Rejected samples and reasons
- Analyzer uptime/downtime and maintenance
- QC pass/fail
- TAT by department and test category
- Reagent/consumable risk
- Historical and live trends

How the Agent Reasons

- Detects bottlenecks before widespread TAT breach
- Identifies departments/analyzers/test categories driving delay

- Correlates downtime, QC or reagent risk with pending workload
- Recognizes recurring delay patterns

Actions the Agent May Perform

- Generate TAT risk alert
- Prioritize operational queue for review
- Recommend workload or process intervention
- Create escalation ticket to lab owner
- Produce daily lab risk summary

When It Escalates

- Critical TAT breach
- STAT delay
- QC failure affecting operations
- Analyzer outage creating capacity risk
- Reagent shortage threatening service continuity

Core KPIs

- TAT breach rate
- Time-to-detect bottleneck
- STAT compliance
- Analyzer availability
- Rejected sample rate
- Intervention effectiveness

Primary Data / Tool Connections

- LIS/LDM
- Analyzer middleware/status feeds
- QC systems
- Inventory/reagent data where available
- Maintenance records

Human / Policy Boundary

No diagnostic or clinical interpretation and no clinical result release. The agent is operational, not clinical.

Legacy RFP basis: SR-004; FR-LAB-001 to FR-LAB-018; FR-EXE-007 and FR-EXE-008.

A4 | MVP

Home Visit Operations Agent

Mission

Improve Home Visit punctuality, field productivity and capacity utilization through live operational intelligence.

What the Agent Monitors

- Scheduled, completed and delayed visits
- Technician/vehicle location where GPS is available
- ETA
- Technician productivity
- Idle time and travel time

- Revenue by technician/route/area/service where available
- Customer satisfaction
- Repeated delays and bottlenecks

How the Agent Reasons

- Predicts delay based on route, location and workload
- Identifies idle capacity or overloaded areas
- Detects repeated technician/route exceptions
- Compares productivity across zones

Actions the Agent May Perform

- Generate delay alert
- Recommend route/capacity rebalancing
- Escalate repeated delays
- Produce area-level capacity view
- Track recovery after intervention

When It Escalates

- High-risk delay clusters
- Repeated missed SLA
- Insufficient field capacity
- Systemic route bottleneck
- Customer-impacting exception requiring manager action

Core KPIs

- On-time arrival rate
- Average delay
- Technician productivity
- Travel vs productive time
- Capacity utilization
- Repeat delay rate

Primary Data / Tool Connections

- Home Visit scheduling
- GPS/fleet
- CRM/order platform
- Payment/revenue data
- CX feedback

Human / Policy Boundary

Automated route recommendations are allowed; dispatch changes should follow approved operating rules or human approval during early phases.

Legacy RFP basis: SR-003; FR-HV-001 to FR-HV-015; FR-EXE-006; IR-006.

A5 | MVP

Customer Experience & Call Center Agent

Mission

Detect service-recovery opportunities, unresolved customer issues and recurring experience problems across branches and channels.

What the Agent Monitors

- Customer complaints and unresolved complaints
- Customer satisfaction metrics and trends
- Call center operational indicators where available
- Customer feedback platform signals
- Waiting-time-related CX risk
- Branch-level complaint patterns

How the Agent Reasons

- Clusters recurring complaint themes
- Links complaints to branch/service/operational cause
- Prioritizes unresolved cases by severity and customer impact
- Identifies contact-center or service bottlenecks

Actions the Agent May Perform

- Create case/ticket and assign owner
- Recommend recovery action
- Escalate overdue customer issues
- Generate branch/CX risk trend
- Feed recurring issues to Quality and Branch agents

When It Escalates

- High-severity complaint
- Repeated unresolved customer issue
- Reputation-sensitive pattern
- Customer safety/quality concern
- Service breakdown requiring executive attention

Core KPIs

- First response time
- Resolution time
- Repeat complaint rate
- Customer satisfaction trend
- Escalation accuracy

Primary Data / Tool Connections

- Call center platform
- CRM
- Customer feedback/CX platform
- Branch complaints
- Quality Hub

Human / Policy Boundary

Sensitive customer communications require approved templates and human authorization unless Wareed explicitly pre-approves an automated workflow.

Legacy RFP basis: SR-001, SR-002; FR-OPS-006; FR-QL-002 and FR-QL-010; IR-009 and IR-010; FR-EXE-010.

Marketing & Growth Intelligence Agent

Mission

Treat marketing as a direct operating and revenue function by linking campaigns, acquisition, offers, products and branch performance to measurable commercial outcomes.

What the Agent Monitors

- Branch revenue against commercial targets
- Campaign performance
- Product and package uptake
- Customer acquisition indicators
- SMS response
- Push notification response
- Call center campaign response
- Digital advertising and promotional offers
- Regional/branch campaign comparisons
- New product/service launch performance

How the Agent Reasons

- Identifies campaign/channel effectiveness by region and branch
- Detects revenue decline and likely commercial drivers
- Compares offers and product uptake
- Identifies weak branches that need activation
- Measures whether corrective commercial actions improved outcomes

Actions the Agent May Perform

- Recommend Stop / Continue / Adjust for campaigns
- Recommend branch-level activation
- Recommend channel/budget reallocation for human approval
- Open marketing action ticket
- Produce campaign ROI and growth opportunity brief

When It Escalates

- Material revenue decline
- Campaign spend with poor outcome
- Offer creating potential margin/value concern where finance data is available
- Major launch underperformance
- Region-wide acquisition decline

Core KPIs

- Incremental revenue
- Conversion rate
- Customer acquisition trend
- Campaign ROI where cost data is integrated
- Offer/product uptake
- Revenue recovery after action

Primary Data / Tool Connections

- CRM
- Marketing automation

- SMS/push platforms
- Call center campaign data
- Digital advertising data
- Branch revenue/finance
- Product/package catalog

Human / Policy Boundary

The agent recommends and prepares actions. It must not launch campaigns, change prices, allocate major budget or message customers without approved authority.

Legacy RFP basis: SR-005; FR-GR-001 to FR-GR-012; FR-EXE-012. Marketing is represented in the legacy RFP primarily through the Growth Hub.

A7 | Extended

Logistics & Courier Agent

Mission

Provide operational visibility and exception management for courier, transport and fleet-supported workflows.

What the Agent Monitors

- Courier/logistics operational status where data is available
- GPS/fleet location where integrated
- Transport delays/exceptions
- Interface health for logistics data

How the Agent Reasons

- Identifies transport exceptions that may affect branch or laboratory flow
- Detects repeated route/service issues
- Correlates courier delay with downstream sample or service impact where data supports it

Actions the Agent May Perform

- Create logistics exception alert
- Assign logistics owner
- Recommend route/service review
- Track closure and recurrence

When It Escalates

- Transport issue threatening sample/service continuity
- Repeated route failure
- Fleet/system outage affecting visibility

Core KPIs

- On-time transport where data is available
- Exception closure time
- Repeat route issue rate
- Downstream impact avoided

Primary Data / Tool Connections

- GPS/fleet management
- Courier/logistics systems

- Operations Hub
- Lab/branch events where correlated

Human / Policy Boundary

The original RFP contains limited direct courier functional requirements; this agent should remain narrow until Wareed defines detailed courier KPIs and operating rules.

Legacy RFP basis: BR-004; FR-OPS-005; IR-006.

A8 | Extended

Product Launch & Commercial Activation Agent

Mission

Ensure every new product or service launch is operationally ready across branches and that early adoption is measured.

What the Agent Monitors

- Launch readiness by branch
- Receipt of launch materials
- Required training completion
- Staff acknowledgment/readiness
- Checklist completion
- Testing/verification/approval
- Planned vs actual launch date
- Early-stage performance and adoption

How the Agent Reasons

- Identifies branches not ready to launch
- Detects missing prerequisites and training gaps
- Correlates readiness with early launch performance
- Prioritizes delayed locations

Actions the Agent May Perform

- Generate readiness alert
- Create prerequisite tasks
- Assign responsible owner
- Escalate delayed readiness
- Produce enterprise adoption report

When It Escalates

- Launch date at risk
- Critical prerequisite missing
- Low adoption after launch
- Multiple branches not ready

Core KPIs

- On-time launch readiness
- Training completion
- Branch readiness rate
- Early adoption
- Time to resolve launch blockers

Primary Data / Tool Connections

- Product Launch Center/work management
- Training/LMS where available
- Marketing/product data
- Branch readiness inputs
- Sales/revenue after launch

Human / Policy Boundary

Does not approve clinical product content or commercial pricing; it manages readiness and adoption workflow.

Legacy RFP basis: FR-PLC-001 to FR-PLC-012; FR-GR-012.

A9 | Extended

Quality & Compliance Agent

Mission

Continuously monitor enterprise quality signals, identify recurring non-compliance and ensure corrective/preventive actions are owned and closed.

What the Agent Monitors

- Customer complaints and trends
- Branch audit results/compliance scores
- Mystery Shopper results where available
- AI camera quality events
- Hygiene alerts
- Uniform/PPE compliance
- Cleaning issues
- Waiting-time-related quality risk
- CAPA status

How the Agent Reasons

- Detects recurring violations
- Ranks quality issues by risk and recurrence
- Correlates quality trends by branch, region, category and owner
- Identifies CAPA overdue or ineffective patterns

Actions the Agent May Perform

- Create quality issue/CAPA task
- Assign owner
- Escalate repeated non-compliance
- Track closure evidence
- Produce trend and hotspot analysis

When It Escalates

- Repeated quality violation
- Critical hygiene/PPE event
- CAPA overdue beyond threshold
- Quality pattern with patient/customer risk

Core KPIs

- CAPA closure time
- Repeat non-compliance rate
- Audit score trend
- Quality alert precision
- Issue recurrence after closure

Primary Data / Tool Connections

- Quality/audit systems
- Mystery Shopper data
- CX/complaints
- AI vision event service
- Branch operating data

Human / Policy Boundary

Supports quality governance; final regulatory/clinical judgments remain with authorized Wareed quality and medical leadership.

Legacy RFP basis: SR-006; FR-QL-001 to FR-QL-014; FR-EXE-009.

A10 | Extended

Workforce & Attendance Agent

Mission

Identify staffing availability risks that affect branch and operational service continuity.

What the Agent Monitors

- Staff attendance
- Workforce availability
- Branch staffing status
- Attendance-related service risk

How the Agent Reasons

- Links staffing gaps to branch wait time or service availability
- Identifies recurring attendance-related risk by location
- Prioritizes operationally material gaps

Actions the Agent May Perform

- Create staffing risk alert
- Notify/assign authorized operations or HR owner
- Recommend temporary operational mitigation for approval
- Track closure

When It Escalates

- Branch unable to operate normally due to staffing
- Repeated critical staffing shortage
- Attendance system/data outage affecting visibility

Core KPIs

- Staffing exception detection time
- Uncovered critical shifts
- Branch service impact from staffing
- Resolution time

Primary Data / Tool Connections

- HR/attendance systems
- Branch Operations Agent
- Operations Hub

Human / Policy Boundary

The original RFP does not define full HR management. This agent is limited to attendance/workforce availability and operational impact; payroll, performance management and HR decisions are out of scope unless separately approved.

Legacy RFP basis: BR-004; SR-001; FR-BR-007; FR-EXE-011; IR-005.

A11 | Extended

Voice of Branch / Field Intelligence Agent

Mission

Turn branch employee feedback into a structured operational intelligence stream rather than unstructured messages.

What the Agent Monitors

- Improvement ideas
- Operational issues/service disruptions
- Administrative/service complaints
- Competitor observations and market intelligence
- Equipment faults
- Training requests
- Supply shortages
- Photos/supporting evidence

How the Agent Reasons

- Automatically classifies submissions
- Prioritizes by severity, urgency and business impact
- Routes to correct owner
- Identifies recurring themes and field hotspots

Actions the Agent May Perform

- Categorize and prioritize submission
- Assign owner
- Track lifecycle to closure
- Generate trend analytics
- Feed competitor intelligence to Marketing/Growth Agent

When It Escalates

- High-impact operational disruption
- Repeated unresolved field issue

- Major competitor signal
- Supply shortage affecting service
- Equipment failure requiring urgent action

Core KPIs

- Time to classification
- Time to owner assignment
- Closure time
- Repeat issue rate
- Actionable idea conversion

Primary Data / Tool Connections

- Voice of Branch portal/app
- Photos/evidence
- Branch user submissions
- Ticketing/work management

Human / Policy Boundary

AI classification and prioritization are allowed; material business decisions remain with accountable owners.

Legacy RFP basis: FR-VOB-001 to FR-VOB-015; SR-008.

A12 | Foundation / Extended

IT & Integration Health Agent

Mission

Protect the reliability of the agentic operating model by continuously monitoring systems, interfaces, data refresh and integration failures.

What the Agent Monitors

- System health
- Integration status
- Data refresh status
- API failures
- Dashboard/agent service availability
- Interface logs and errors
- Retry/failure-recovery status
- Security-relevant operational events where integrated

How the Agent Reasons

- Distinguishes source-system issue from integration issue
- Detects stale data before business agents act on it
- Identifies recurring interface failure
- Prioritizes failures by business impact

Actions the Agent May Perform

- Create IT incident
- Trigger approved automated retry
- Escalate persistent failures
- Mark affected data/agent output as degraded

- Produce integration health report

When It Escalates

- Critical integration outage
- Stale/invalid data affecting decision agents
- Repeated API failure
- Security or availability issue with enterprise impact

Core KPIs

- Integration uptime
- API failure rate
- Mean time to recover
- Data freshness compliance
- Automated retry success

Primary Data / Tool Connections

- API gateway
- Middleware/integration platform
- Monitoring/logging
- Source-system health
- Data pipeline telemetry

Human / Policy Boundary

May execute only pre-approved technical recovery actions. Security-sensitive or production-impacting changes require authorized IT approval.

Legacy RFP basis: SR-007; IR-001 to IR-018; NFR-001 to NFR-012; CSR controls as applicable.

S1 | Foundation

Data Quality & Governance Agentic Service

Mission

Ensure that every business agent operates on trusted, governed and auditable data.

What the Agent Monitors

- Data completeness, accuracy, consistency and integrity
- Refresh frequency/synchronization
- Master data consistency
- Data quality rule violations
- Retention and auditability

How the Agent Reasons

- Detects data anomalies that can invalidate agent conclusions
- Separates missing data from real operational exception
- Tracks data ownership and unresolved quality issues

Actions the Agent May Perform

- Flag unreliable datasets
- Quarantine or mark degraded data for agent use
- Create data-quality issue

- Produce quality score by domain
- Maintain audit evidence

When It Escalates

- Critical data-quality failure
- Stale data supporting high-risk decision
- Master data conflict
- Repeated data ownership gap

Core KPIs

- Data quality score
- Freshness compliance
- Issue resolution time
- False operational alerts caused by data quality

Primary Data / Tool Connections

- Operational data platform
- MDM
- Data pipelines
- Data catalog/governance metadata

Human / Policy Boundary

Does not alter authoritative source data without approved data stewardship workflow.

Legacy RFP basis: DR-001 to DR-012; IR-013 to IR-015.

S2 | Foundation

Alert, Ticketing & Corrective Action Agentic Service

Mission

Provide one governed mechanism for turning operational exceptions into owned, timed and auditable resolution.

What the Agent Monitors

- All alert rules
- Severity levels
- Owner assignment
- Response targets
- Acknowledgment
- Escalation
- Closure evidence
- Alert trends and recurrence

How the Agent Reasons

- Deduplicates and suppresses low-value alerts
- Prioritizes by Information / Warning / Critical / Executive Escalation
- Identifies overdue actions
- Detects recurring issue patterns

Actions the Agent May Perform

- Create alert/ticket

- Assign owner
- Start SLA timer
- Escalate automatically
- Close only after controlled resolution
- Produce alert analytics

When It Escalates

- SLA breach
- Executive severity
- Repeated unresolved alert
- No accountable owner
- Failed corrective action

Core KPIs

- Alert fatigue reduction
- Acknowledgment time
- Resolution time
- SLA compliance
- Repeat alert frequency

Primary Data / Tool Connections

- All agents
- Ticketing/work management
- Email/SMS/notification systems

Human / Policy Boundary

Closure rules must be controlled and auditable; executive escalation cannot be suppressed by an agent without approved policy.

Legacy RFP basis: AR-001 to AR-012; BR-007 and BR-010.

S3 | Extended

AI Vision Event Agentic Service

Mission

Convert approved camera analytics into event-driven operational signals instead of passive surveillance.

What the Agent Monitors

- No receptionist event where feasible
- Crowding/queue congestion
- Cleaning/hygiene events
- PPE compliance
- Uniform compliance
- Unattended service desk
- Camera offline/communication failure

How the Agent Reasons

- Classifies camera event type
- Scores confidence and business relevance
- Suppresses repeated false-positive patterns according to approved rules

- Links events to branch/quality workflow

Actions the Agent May Perform

- Create event for Branch or Quality Agent
- Attach evidence metadata where permitted
- Track camera health
- Escalate repeated high-confidence violations

When It Escalates

- Repeated quality/compliance event
- Critical service-desk absence
- Persistent queue congestion
- Camera service outage where coverage is required

Core KPIs

- Detection precision/recall where validated
- False-positive rate
- Event-to-action time
- Repeat violation rate

Primary Data / Tool Connections

- AI camera/video analytics
- VMS
- Branch metadata
- Quality Hub

Human / Policy Boundary

Privacy-aware, event-driven only. No facial identification or unsupported surveillance use case is included by this proposal. Vendor must define limitations, accuracy and privacy controls.

Legacy RFP basis: AIR-CAM-001 to AIR-CAM-012; CSR-014; IR-007 and IR-008.

6. Human-in-the-Loop Authority Model

Authority Level	AI Capability	Governance
Level 1 - Observe	Read data, summarize status, detect anomalies.	No approval required for read-only analysis.
Level 2 - Recommend	Generate root-cause hypotheses and corrective recommendations.	Human decides whether to act.
Level 3 - Assign	Create approved ticket/task and route to defined owner.	Allowed only under approved workflow rules.
Level 4 - Escalate	Escalate overdue, repeated or high-severity issues.	Automatic within approved escalation policy.
Level 5 - Act with Approval	Execute low-risk workflow action after explicit approval or under pre-approved policy.	Audited and reversible where possible.

Explicit High-Risk Boundaries

- No clinical diagnosis, interpretation or result release.
- No unapproved pricing changes or discount activation.
- No autonomous campaign launch or major budget reallocation.
- No sensitive customer communication outside approved templates/policies.
- No HR employment decision based solely on AI output.
- No privileged production/security change without authorized IT approval.

7. Sovereign AI, Data and Cybersecurity

Sovereign AI is a Wareed strategic requirement layered on top of the legacy RFP security controls. The selected bidder must propose an architecture that keeps Wareed in control of its data, agent policies, audit evidence, integrations and operational knowledge.

Mandatory Design Expectations

- Approved Saudi-based or sovereign deployment model for sensitive workloads, subject to Wareed legal/cybersecurity approval.
- Wareed ownership of business, operational, customer, laboratory, financial, employee, branch and performance data.
- Role-Based Access Control, enterprise authentication, least privilege and privileged-access controls.
- Encryption in transit and at rest; secure APIs; network segmentation where applicable.
- Complete audit logging for user, administrative, agent and security-relevant actions.
- Agent-level tool permissions, policy boundaries and approval gates.
- No use of Wareed data to train third-party/shared models without explicit written approval.
- Backup, recovery, disaster recovery, patching, vulnerability management and support SLAs.
- Privacy-aware AI camera controls and documented accuracy / false-positive management.

Data Ownership

The legacy procurement appendix explicitly states that Wareed business, operational, customer, laboratory, financial, employee, branch and performance data remain Wareed property. The Agentic AI contract should extend this principle to prompts, agent traces, business rules, embeddings, derived operational intelligence and custom agent configurations where contractually feasible.

8. Integration and Data Foundation

The agentic platform cannot outperform the trustworthiness and timeliness of its data. Phase 1 and Phase 2 therefore prioritize integrations, data quality, identity, auditability and operational event architecture before broad autonomy.

Required Integration Domains from the Legacy RFP

- Laboratory Information System (LIS/LDM)
- ERP / finance systems
- CRM
- HR and attendance
- GPS / fleet management
- CCTV / VMS
- AI camera / video analytics
- Call center
- Customer feedback / CX platforms
- Standards-based APIs and middleware

Integration Requirements

- Central interface health monitoring and data freshness status.
- Centralized logging, auditing and error reporting.
- Automated retry, exception handling and failure recovery.
- Integration Responsibility Matrix defining Wareed vs bidder responsibilities.
- Explicit assumptions about source-system readiness and API/access availability.
- Enterprise data model, master data principles, retention and data-quality rules.

9. Phased Implementation Roadmap

Phase 1 - Discovery, Sovereign Architecture & Agent Design

Indicative duration: 4-8 weeks

- Current-state and source-system assessment
- Agent catalogue confirmation and use-case prioritization
- Sovereign hosting/security decision
- Data and integration readiness
- Human authority / approval matrix
- MVP definition and baseline KPIs

Phase 2 - Data, Integration & Agentic Control Foundation

Indicative duration: 8-16 weeks

- Operational data layer and integrations
- Identity, RBAC, audit and security controls
- Data Quality & Governance service
- Alert/Ticketing/Corrective Action service
- Agentic Orchestrator
- Observability, logging, testing framework

Phase 3 - AI-Assisted Operations MVP

Indicative duration: 8-12 weeks

- Executive Intelligence & CEO Briefing Agent
- Branch Operations Agent
- Main Laboratory Operations & TAT Agent
- Home Visit Operations Agent
- Customer Experience & Call Center Agent
- Marketing & Growth Intelligence Agent

Phase 4 - Extended Agents & Semi-Autonomous Workflows

Indicative duration: 12-24 weeks

- Quality & Compliance Agent
- IT & Integration Health Agent
- Logistics & Courier Agent
- Product Launch Agent
- Workforce & Attendance Agent
- Voice of Branch Agent
- AI Vision Event service
- Approved task assignment, SLA and escalation automation

Phase 5 - Predictive & Sovereign Scale

Indicative duration: Ongoing

- Predictive branch decline
- TAT risk prediction
- Home Visit capacity forecasting

- Marketing/commercial performance prediction
- Quality recurrence prediction
- Cross-agent planning and executive scenario analysis
- Expansion to new functions and services

10. Proof of Concept

Shortlisted bidders should be required to demonstrate agentic behavior, not only dashboards. The POC should use representative or synthetic Wareed scenarios and show traceable evidence from signal to recommendation to governed action.

POC Scenario	Expected Demonstration
CEO Daily Brief	Generate a concise brief from multi-domain sample data and identify three decisions requiring executive attention.
Branch Agent	Detect a branch decline and explain the combined revenue, waiting-time, staffing and complaint signals.
Lab Agent	Detect TAT risk, identify likely bottleneck and recommend an operational response without making a clinical decision.
Home Visit Agent	Predict/identify delay, recommend route or capacity action and escalate repeated delay.
Marketing Agent	Compare campaigns/branches, recommend Stop / Continue / Adjust and show the business evidence.
Alert-to-Ticket	Create an owned corrective action, start SLA timer, escalate overdue item and preserve audit history.
Integration Health	Simulate an API/data-refresh failure and show that affected agent outputs are flagged as degraded.
AI Governance	Show approval workflow, tool permissions, action logs and human override.
Vision Event - Optional	Demonstrate event-driven quality event with privacy and confidence controls.

11. Acceptance and Success Measures

Technical Acceptance

- End-to-end system integration testing (SIT).
- User acceptance testing (UAT).
- Validation of KPIs, dashboards and agent outputs.
- Validation of alerts, escalation and ticket workflows.
- Performance/scalability testing.
- Cybersecurity verification and vulnerability testing.
- Backup, failover and recovery testing where applicable.
- Operational Readiness Testing (ORT).

Agentic Acceptance

- Agent recommendation accuracy measured against agreed test cases.
- No unauthorized tool/action execution.
- Every action/recommendation has traceable evidence and audit history.
- Data freshness and data-quality state visible to agents and users.
- Escalation policies operate exactly as approved.
- Human override works at all times.
- False-positive/false-negative thresholds are documented for AI vision where used.
- Business KPIs improve or show a credible baseline-to-target path during pilot.

Business Success Measures

- Reduction in manual follow-up effort.
- Faster issue detection and owner assignment.
- Improved SLA/corrective-action closure.
- Reduced TAT/branch/Home Visit exceptions in targeted pilots.
- Improved marketing decision speed and campaign effectiveness.
- Lower alert fatigue with better prioritization.
- Higher executive situational awareness with fewer raw reports.

12. Vendor Response Requirements

The bidder should respond with a complete strategic and technical proposal. The response should preserve the structured, auditable approach of the original procurement model but replace physical-command-center emphasis with agentic AI architecture and sovereign operation.

1. Executive proposal and understanding of Wareed
2. Agentic operating model and agent catalogue
3. Sovereign AI / hosting architecture
4. Solution architecture
5. Integration architecture and Responsibility Matrix
6. Data architecture, governance and quality model
7. Agentic control plane, policy and approval model
8. Detailed role definition for each proposed agent
9. Dashboard / executive experience concept
10. Alert, ticketing and corrective-action workflow
11. Cybersecurity and privacy architecture
12. POC plan
13. Phased implementation schedule
14. Testing and acceptance plan
15. Training and change-management plan
16. Support / SLA and managed service model
17. Commercial proposal and recurring-cost model
18. Risk register
19. Assumptions, exclusions and deviation log
20. IP, data ownership, model ownership and portability position

13. Recommended Evaluation Model

The legacy RFP evaluation included physical command center design and AV/video wall solution. For this AI-first procurement, those factors should be removed from mandatory technical weighting and treated as optional value-added scope. The following weighting is recommended by Wareed for the AI-company engagement.

Evaluation Area	Weight
Business & healthcare understanding	10%
Agentic AI architecture and practical autonomy model	20%
Integration and data architecture	20%
Sovereign AI, cybersecurity and governance	15%
Operating model and workflow design	10%
POC quality and measurable value	10%
Implementation methodology and change management	5%
Support, SLA and local/regional capability	5%
Commercial value / TCO	5%

Pass / Fail Recommendation

Demonstrated integration capability, healthcare relevance, AI governance, data security, auditable agent actions, implementation schedule, local/regional support, and a successful POC should be mandatory pass/fail conditions.

14. Commercial Structure

Wareed should request a phased commercial proposal with measurable milestone acceptance. Physical room, video wall, AV and operator furniture should be separated as optional items and must not obscure the recurring economics of the AI platform.

- Professional services - discovery, operating model, agent design and business analysis.
- Sovereign platform / cloud / infrastructure costs.
- Agentic AI platform and model usage/licensing.
- Data integration and API/connectors.
- Data platform and governance components.
- Cybersecurity and observability.
- Agent development and testing by phase.
- Dashboard/mobile executive experience.
- Training and change management.
- Go-live and stabilization.
- Support, managed service and SLA.
- Future agent expansion and continuous improvement.
- Optional physical command room / displays / AV / furniture.

Payment Milestone Logic

- Discovery accepted
- Architecture and operating model approved
- Integration/data foundation accepted
- MVP agents pass UAT/ORT
- POC/business pilot targets achieved
- Extended automation accepted
- Go-live and stabilization
- Final acceptance

Appendix A - RFP-to-Agent Mapping

Legacy RFP Domain	Requirement Basis	Agentic Replacement
Business coordination / enterprise visibility	BR-001 to BR-010	A0 Orchestrator, A1 CEO Agent
Executive dashboard	FR-EXE-001 to FR-EXE-015	A1 Executive Intelligence & CEO Briefing
Operations Hub	FR-OPS-001 to FR-OPS-010	A0, A2, A4, A7, A5
Branch monitoring	FR-BR-001 to FR-BR-015	A2 Branch Operations
Home Visit	FR-HV-001 to FR-HV-015	A4 Home Visit
Main laboratory	FR-LAB-001 to FR-LAB-018	A3 Main Laboratory
Growth / marketing	FR-GR-001 to FR-GR-012	A6 Marketing & Growth
Quality Hub	FR-QL-001 to FR-QL-014	A9 Quality & Compliance
Intelligence Hub	FR-INT-001 to FR-INT-010	A0 Orchestrator, A1 CEO Agent
Voice of Branch	FR-VOB-001 to FR-VOB-015	A11 Voice of Branch
Product Launch Center	FR-PLC-001 to FR-PLC-012	A8 Product Launch
AI camera events	AIR-CAM-001 to AIR-CAM-012	S3 AI Vision Event Service
Integrations	IR-001 to IR-018	A12 IT & Integration Health
Data	DR-001 to DR-012	S1 Data Quality & Governance
Alerts	AR-001 to AR-012	S2 Alert/Ticketing/Corrective Action
Non-functional / continuity	NFR-001 to NFR-012	A12 + platform architecture
Cybersecurity	CSR-001 to CSR-014	Platform security + all agents
Workforce attendance	SR-001, FR-BR-007, FR-EXE-011, IR-005	A10 Workforce & Attendance
Courier/logistics	BR-004, FR-OPS-005, IR-006	A7 Logistics & Courier

Closing Statement

Wareed is not seeking a dashboard vendor, a conventional control-room integrator, or a generic AI chatbot. Wareed is seeking a strategic AI transformation partner capable of building a sovereign, secure, governed and scalable agentic operating system that converts enterprise data into continuous operational action and executive decision intelligence.

Requested Next Step

The AI company should return a proposed architecture, detailed agent catalogue, sovereignty model, phased roadmap, POC plan, commercial structure, risks and assumptions. Wareed will then shortlist the solution based on measurable business value and implementation credibility.